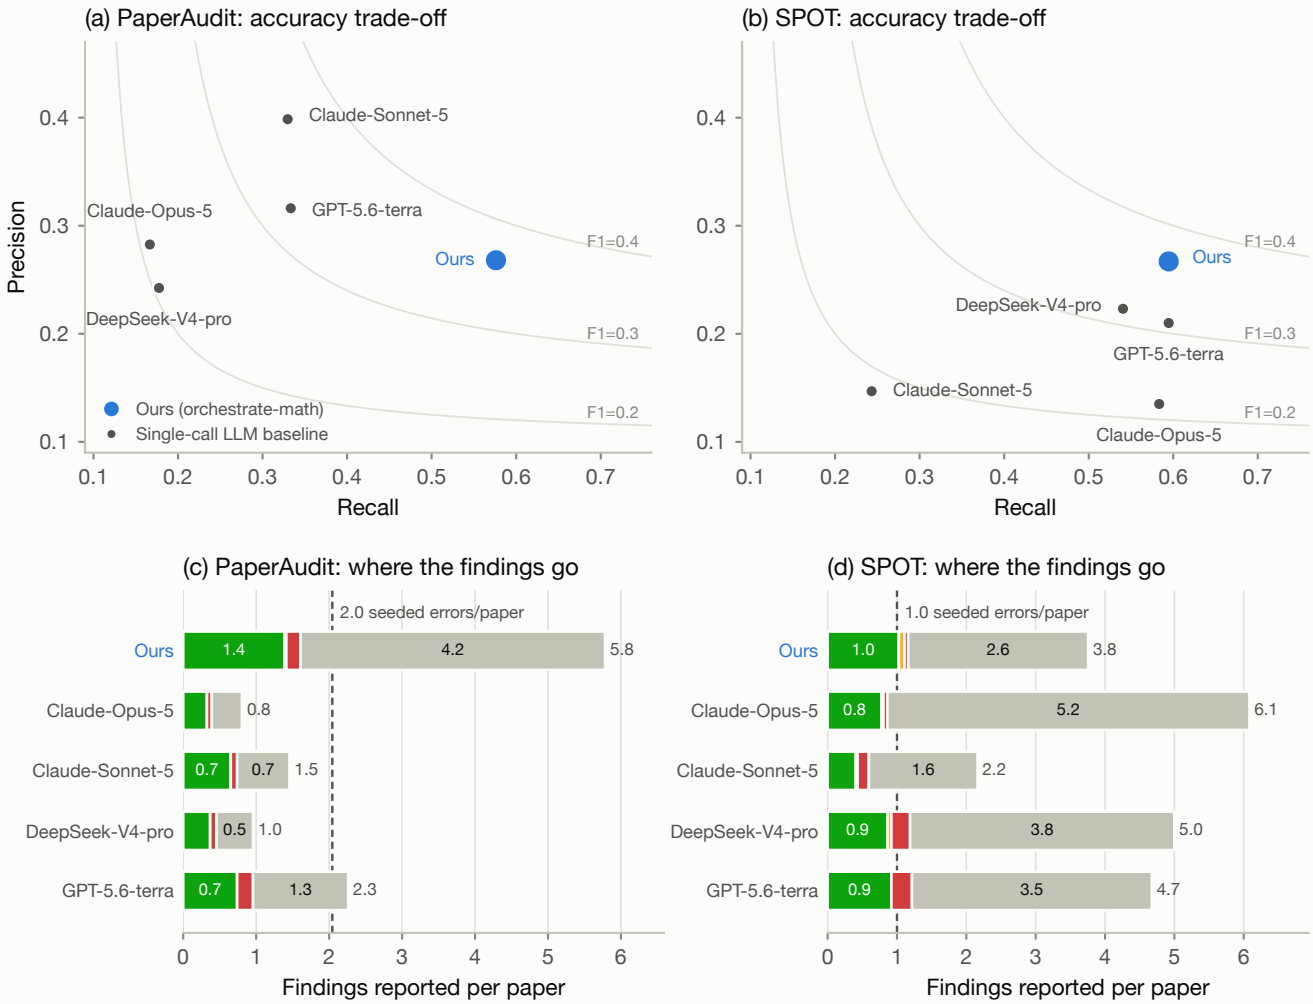
(a) PaperAudit: accuracy trade-off
Precision
Recall
Claude-Sonnet-5
GPT-5.6-terra
Claude-Opus-5
Ours
DeepSeek-V4-pro
F1=0.4
F1=0.3
F1=0.2
Ours (orchestrate-math)
Single-call LLM baseline
(b) SPOT: accuracy trade-off
Recall
Ours
DeepSeek-V4-pro
GPT-5.6-terra
Claude-Sonnet-5
Claude-Opus-5
F1=0.4
F1=0.3
F1=0.2
(c) PaperAudit: where the findings go
2.0 seeded errors/paper
Ours 1.4 4.2 5.8
Claude-Opus-5 0.8
Claude-Sonnet-5 0.7 0.7 1.5
DeepSeek-V4-pro 0.5 1.0
GPT-5.6-terra 0.7 1.3 2.3
Findings reported per paper
(d) SPOT: where the findings go
1.0 seeded errors/paper
Ours 1.0 2.6 3.8
Claude-Opus-5 0.8 5.2 6.1
Claude-Sonnet-5 1.6 2.2
DeepSeek-V4-pro 0.9 3.8 5.0
GPT-5.6-terra 0.9 3.5 4.7
Findings reported per paper
true positive (credited)
duplicate
spurious (wrong)
out of scope (true, but not a seeded error)